

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: # import plotly modules
import plotly.offline as py
py.init_notebook_mode(connected=True)
import plotly.graph_objs as go
import plotly.tools as tls

import warnings
warnings.filterwarnings('ignore')

import nltk
nltk.download()
#nltk.download('punkt')
#nltk.download('stopwords')
from nltk import word_tokenize, sent_tokenize
from nltk import pos_tag
from nltk.corpus import stopwords
```

showing info https://raw.githubusercontent.com/nltk/nltk_data/gh-pages/index.xml
1 (https://raw.githubusercontent.com/nltk/nltk_data/gh-pages/index.xml)

```
In [3]: hotel_review = pd.read_csv('hotelreview.csv', low_memory=False)
```

```
In [4]: hotel_review.shape
```

```
Out[4]: (35912, 19)
```

```
In [5]: hotel_review.head()
```

```
Out[5]:
```

	address	categories	city	country	latitude	longitude	name	postalCode	province	rev
0	Riviera San Nicol 11/a	Hotels	Mableton	US	45.421611	12.376187	Hotel Russo Palace	30126	GA	22T
1	Riviera San Nicol 11/a	Hotels	Mableton	US	45.421611	12.376187	Hotel Russo Palace	30126	GA	03T
2	Riviera San Nicol 11/a	Hotels	Mableton	US	45.421611	12.376187	Hotel Russo Palace	30126	GA	13T
3	Riviera San Nicol 11/a	Hotels	Mableton	US	45.421611	12.376187	Hotel Russo Palace	30126	GA	27T
4	Riviera San Nicol 11/a	Hotels	Mableton	US	45.421611	12.376187	Hotel Russo Palace	30126	GA	05T

```
In [6]: eighty_count = len(hotel_review)*4 / 5
```

```
In [7]: hotel_review1 = hotel_review.dropna(thresh=eighty_count,axis=1)
```

```
In [8]: hotel_review1.shape
```

```
Out[8]: (35912, 15)
```

```
In [9]: object_columns_df = hotel_review1.select_dtypes(include=['object'])
print(object_columns_df.iloc[0])

cols = ['country']
for name in cols:
    print(name, ':')
    print(object_columns_df[name].value_counts(), '\n')
```

```
address                Riviera San Nicol 11/a
categories              Hotels
city                   Mableton
country                US
name                   Hotel Russo Palace
postalCode             30126
province               GA
reviews.date            2013-09-22T00:00:00Z
reviews.dateAdded      2016-10-24T00:00:25Z
reviews.text            Pleasant 10 min walk along the sea front to th...
reviews.title           Good location away from the crouds
reviews.username        Russ (kent)
Name: 0, dtype: object
country :
US      35912
Name: country, dtype: int64
```

```
In [10]: old_names = ['reviews.date', 'reviews.rating', 'reviews.title', 'reviews.text']
new_names = ['date', 'rating', 'title', 'text']
hotel_review2 = hotel_review1.rename(columns=dict(list(zip(old_names, new_names))))
```

```
In [11]: df = hotel_review2[['latitude', 'longitude', 'name', 'address', 'postalCode', 'ca'
```

In [12]: `df.head()`

Out[12]:

	latitude	longitude	name	address	postalCode	categories	city	country	date
0	45.421611	12.376187	Hotel Russo Palace	Riviera San Nicol 11/a	30126	Hotels	Mableton	US	2013-02-22T00:00:00
1	45.421611	12.376187	Hotel Russo Palace	Riviera San Nicol 11/a	30126	Hotels	Mableton	US	2015-03-03T00:00:00
			Hotel	Riviera					2014-0

In [13]: `df.to_csv("cleaned_hotelreview.csv", index=False)`

In [14]: `df = df[pd.notnull(df['name'])]
df = df[pd.notnull(df['latitude'])]
df = df[pd.notnull(df['longitude'])]
df = df[pd.notnull(df['rating'])]
df = df[pd.notnull(df['date'])]`

Q1. Which hotel has the highest number of reviews.

In [15]: `q1 = df['name'].value_counts().reset_index().iloc[0]['index']
print("Answer: " + q1)`

Answer: The Alexandrian, Autograph Collection

Q2. Which hotel has the highest average rating of reviews.

In [16]: `q2 = df.groupby('name')['rating'].mean().reset_index().sort_values(by='rating', ascending=False).iloc[0]['name']
print("Answer: " + q2)`

Answer: Pelican Shores Inn

```

In [17]: q2 = df.groupby('name')['rating'].mean().reset_index().sort_values(by='rating', a
trace = go.Bar(
    x=q2['name'],
    y=q2['rating'],
    marker=dict(
        color='rgb(158,202,225)',
        line=dict(
            color='rgb(8,48,107)',
            width=1.5,
        )
    ),
    opacity=0.6
)


data = [trace]

layout = go.Layout(
    title='Bar Chat Showing Top 5 Hotels With Highest Average Ratings.',
)

fig = go.Figure(data=data, layout=layout)
py.iplot(fig, filename='hotel-reviews-highest-rating')

```

Bar Chat Showing Top 5 Hotels With Highest A



Q4. Which City has the highest number of hotels.

```
In [18]: q3 = df['city'].value_counts().reset_index().iloc[0]['index']  
print("Answer: " + q3)
```

Answer: Alexandria


```
In [19]: q3 = df['city'].value_counts()[:10]
trace = go.Bar(
    x=q3.index,
    y=q3.values,
    marker=dict(
        color='rgb(158,202,225)',
        line=dict(
            color='rgb(8,48,107)',
            width=1.5,
        )
    ),
    opacity=0.6
)

data = [trace]

layout = go.Layout(
    title='Bar Chart Showing Top 10 Cities With Highest Reviews.',
)

fig = go.Figure(data=data, layout=layout)
py.iplot(fig, filename='hotel-reviews-highest-cities')
```

Bar Chart Showing Top 10 Cities With High



Q3. What is the relationship between total number of reviews per hotel and average rating of the hotel.


```

In [20]: group_name = df.groupby(['name'])['rating'].mean().reset_index()

group_count = df.groupby('name').count().reset_index()
old_names = ['latitude']
new_names = ['count']
group_count.rename(columns=dict(zip(old_names, new_names)), inplace=True)
group_count = group_count[['name', 'count']]

q4 = pd.merge(group_name, group_count, left_index=True, right_index=True)[['name_',
# q4.plot.scatter(x='rating', y='count')

x = (q4['rating']).values
y = (q4['count']).values

data = go.Data([
    go.Scatter(
        x = x,
        y = y,
        mode = 'markers'
    )
])


layout = go.Layout(
    title='Diagram Showing The Relationship Between Ratings & Number Of Reviews.'
)


fig = dict(data=data, layout=layout)

# Plot and embed in ipython notebook!
py.iplot(fig, filename='hotels-reviews-scatter')

```

Diagram Showing The Relationship Between Ratings





Answer: Based on the above scatter plot diagram, there is no relationship between the average rating of an hotel and the total number of reviews of the hotel.

Q5. Plot an interactive map of the hotels and average review ratings as a label.

```

In [21]: q5 = df.groupby(['name', 'latitude', 'longitude'])['rating'].mean().reset_index()
lat = q5.latitude
lon = q5.longitude
name = q5.name
rating = round(q5.rating,2)

mapbox_access_token = 'pk.eyJ1Ijoia2FtcGFyaWEiLCJhIjoib0JLTExtSSJ9.6ahf835RV3kBUm
data = go.Data([
    go.Scattermapbox(
        lat=lat,
        lon=lon,
        mode='markers',
        marker=go.Marker(
            size=10,
            color='rgb(255, 0, 0)',
            opacity=0.7
        ),
        text=rating,
        hoverinfo=''
    ),
    go.Scattermapbox(
        lat=lat,
        lon=lon,
        mode='markers',
        marker=go.Marker(
            size=8,
            color='rgb(242, 177, 172)',
            opacity=0.7
        ),
        text=rating,
        hoverinfo=''
    )
])


layout = go.Layout(
    title='Interactive Map Showing The Location Of Hotels & Average Ratings.',
    autosize=True,
    hovermode='closest',
    showlegend=False,
    mapbox=dict(
        accesstoken=mapbox_access_token,
        bearing=0,
        center=dict(
            lat=38,
            lon=-94
        ),
        pitch=0,
        zoom=3,
        style='dark'
    ),
)

fig = dict(data=data, layout=layout)

py.iplot(fig, filename='hotel-reviews-map')

```

Interactive Map Showing The Location Of Hote



```
In [22]: # filterout ratings that are zero
hotel_review2.shape
```

```
Out[22]: (35912, 15)
```

```
In [23]: hotel_review.columns
```

```
Out[23]: Index(['address', 'categories', 'city', 'country', 'latitude', 'longitude',
               'name', 'postalCode', 'province', 'reviews.date', 'reviews.dateAdded',
               'reviews.doRecommend', 'reviews.id', 'reviews.rating', 'reviews.text',
               'reviews.title', 'reviews.userCity', 'reviews.username',
               'reviews.userProvince'],
              dtype='object')
```

```
In [24]: hotel_review = hotel_review[hotel_review['reviews.rating']>0]
```

```
In [25]: hotel_review_data = hotel_review.rename(index=str, columns={'reviews.date': 'reviewdate',
                           'reviews.doRecommend': 'reviewsdoRecommend', 'reviews.id': 'reviewsid', 'reviews.rating': 'reviewsrating',
                           'reviews.title': 'reviewstitle', 'reviews.userCity': 'reviewsuserCity', 'reviews.userProvince': 'reviewsuserProvince'})
```

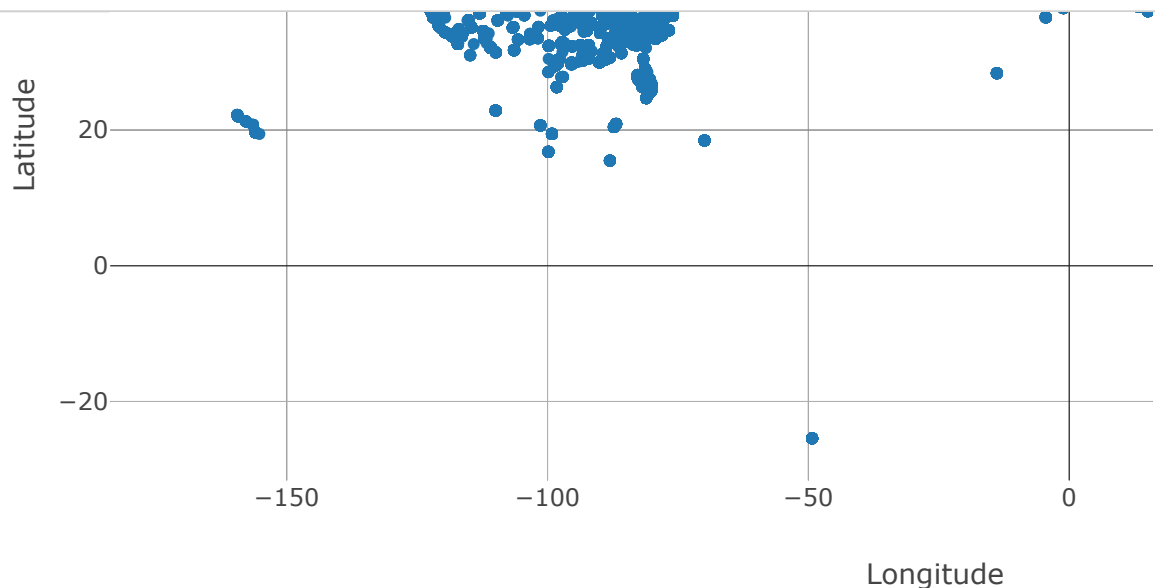
```
In [26]: hotel_review_data.reviewstext = hotel_review_data.reviewstext.fillna('x')
#A few hundred ratings had a score above 5, filtering these out
hotel_review_data = hotel_review_data[hotel_review_data['reviewsrating']<=5]
#A few hundred ratings had decimals, rounding each of those down to an integer
hotel_review_data.reviewsrating = hotel_review_data.reviewsrating.astype(int)
```

```
In [27]: #Creating a function that I will use to clean review strings
#Function makes the string 'txt' lowercase, removes stopwords, finds the length, and
#Returns a list of the length, cleaned txt, and only adjective txt
def cleanme(txt):
    sent = txt.lower()
    wrds = word_tokenize(sent)
    clwrds = [w for w in wrds if not w in stopwords.words('english')]
    ln = len(clwrds)
    pos = pd.DataFrame(pos_tag(wrds))
    pos = " ".join(list(pos[pos[1].str.contains("JJ")].iloc[:,0]))
    rt = [ln, " ".join(clwrds), pos]
    return(rt)
```

```
In [28]: hotel_review_data.country.unique()
```

```
Out[28]: array(['US'], dtype=object)
```

```
In [29]: plt_review = go.Scatter(x = hotel_review_data.longitude, y=hotel_review_data.latitude)
lyt_review = go.Layout(title="Locations of Hotel Reviews", xaxis=dict(title='Longitude',
fig_review = go.Figure(data = [plt_review], layout=lyt_review)
py.iplot(fig_review)
```



```
In [30]: #Create a field that shows the length of each review
review_length = list()
for i in range(len(hotel_review_data)):
    review_length.append(cleanme(hotel_review_data.iloc[i,:]['reviewtext']))
review_length = pd.DataFrame(review_length)
review_length.columns = ['reviewlen', 'cleanrev', 'adjreview']
```

```
In [31]: #Add calculated columns back to the dataset
hotel_review_data = hotel_review_data.reset_index()
hotel_review_data = pd.concat([hotel_review_data,review_length], axis=1)
hotel_review_data.head()
```


Out[31]:

	index	address	categories	city	country	latitude	longitude	name	postalCode	provin
0	0	Riviera San Nicol 11/a	Hotels	Mableton	US	45.421611	12.376187	Hotel Russo Palace	30126	(
1	1	Riviera San Nicol 11/a	Hotels	Mableton	US	45.421611	12.376187	Hotel Russo Palace	30126	(
2	2	Riviera San Nicol 11/a	Hotels	Mableton	US	45.421611	12.376187	Hotel Russo Palace	30126	(
3	3	Riviera San Nicol 11/a	Hotels	Mableton	US	45.421611	12.376187	Hotel Russo Palace	30126	(
4	4	Riviera San Nicol 11/a	Hotels	Mableton	US	45.421611	12.376187	Hotel Russo Palace	30126	(

5 rows × 23 columns


```
In [32]: plt_hist = go.Histogram(x = hotel_review_data.reviewlen)
lyt_hist = go.Layout(title="Frequency of Review Length", xaxis=dict(title='Review
fig_hist = go.Figure(data=[plt_hist], layout=lyt_hist)
py.iplot(fig_hist)
```

Frequency of Review Lengt



```
In [33]: hotel_review_data = hotel_review_data.sort_values(by='reviewlen')
plt_scatter = go.Scatter(x = hotel_review_data.reviewlen, y = hotel_review_data.rating)
lyt_scatter = go.Layout(title="Review Length vs. Star Rating", xaxis=dict(title='Review Length', range=[0, 250]), yaxis=dict(title='Star Rating', range=[1, 5]))
fig_scatter = go.Figure(data=[plt_scatter], layout=lyt_scatter)
py.iplot(fig_scatter)
print("Review Length to Rating Correlation:", hotel_review_data.reviewlen.corr(hotel_review_data.rating))
```

Review Length vs. Star Rating



Review Length to Rating Correlation: -0.149136176154

```
In [34]: #Setting up the X and Y data, where X is the review text and Y is the rating
#Three different inputs will be used: original review text, cleaned review text, and adjusted review text
x1 = hotel_review_data.reviewstext
x2 = hotel_review_data.cleanrev
x3 = hotel_review_data.adjreview
y = hotel_review_data.reviewsrating
```



```
In [35]: from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer

#Creating a vectorizer to split the text into unigrams and bigrams
vect = TfidfVectorizer(ngram_range = (1,2))
x_vect1 = vect.fit_transform(x1)
x_vect2 = vect.fit_transform(x2)
x_vect3 = vect.fit_transform(x3)


In [36]: from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

#Making some simple functions for linear svc, knn, and naive bayes
def linsvc(x,y):
    x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=42)
    classf = LinearSVC()
    classf.fit(x_train, y_train)
    pred = classf.predict(x_test)
    print("Linear SVC:", accuracy_score(y_test, pred))
    return(y_test, pred)

def revknn(x,y):
    x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=42)
    classf = KNeighborsClassifier(n_neighbors=2)
    classf.fit(x_train, y_train)
    pred = classf.predict(x_test)
    print("kNN:", accuracy_score(y_test, pred))
    return(y_test, pred)

def revnb(x,y):
    x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=42)
    classf = MultinomialNB()
    classf.fit(x_train, y_train)
    pred = classf.predict(x_test)
    print("Naive Bayes:", accuracy_score(y_test, pred))
    return(y_test, pred)
```

```
In [37]: from sklearn.svm import LinearSVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import MultinomialNB, GaussianNB

svmy1,svmp1 = linsvc(x_vect1,y)
svmy2,svmp2 = linsvc(x_vect2,y)
svmy3,svmp3 = linsvc(x_vect3,y)

knny1,knnp1 = revknn(x_vect1,y)
knny2,knnp2 = revknn(x_vect2,y)
knny3,knnp3 = revknn(x_vect3,y)

nby1,nbp1 = revnb(x_vect1,y)
nby2,nbp2 = revnb(x_vect2,y)
nby3,nbp3 = revnb(x_vect3,y)
```

```
Linear SVC: 0.522295959949
Linear SVC: 0.519734544184
Linear SVC: 0.451042030504
kNN: 0.300034928397
kNN: 0.300500640354
kNN: 0.298987076493
Naive Bayes: 0.403190126907
Naive Bayes: 0.406566538596
Naive Bayes: 0.421236465246
```

```
In [38]: #This function will plot a confusion matrix and is taken from the sklearn document
def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = np.around((cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]),decim
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```

In [39]: `import itertools`

```
c1 = confusion_matrix(svm1,svmp1)
c2 = confusion_matrix(svm2,svmp2)
c3 = confusion_matrix(nby2,nbp2)
class_names = ['1', '2', '3', '4', '5']
plt.figure()
plot_confusion_matrix(c1, classes=class_names,normalize=False,title='Confusion ma
plt.figure()
plot_confusion_matrix(c2, classes=class_names,normalize=False,title='Confusion ma
plt.figure()
plot_confusion_matrix(c3, classes=class_names,normalize=False,title='Confusion ma
```

Confusion matrix, without normalization


```
[[ 443   89   75   71   72]
 [ 174  152  239  115   78]
 [ 108   89  449  543  278]
 [   28   43  277 1083 1025]
 [   28   11   80  680 2359]]
```

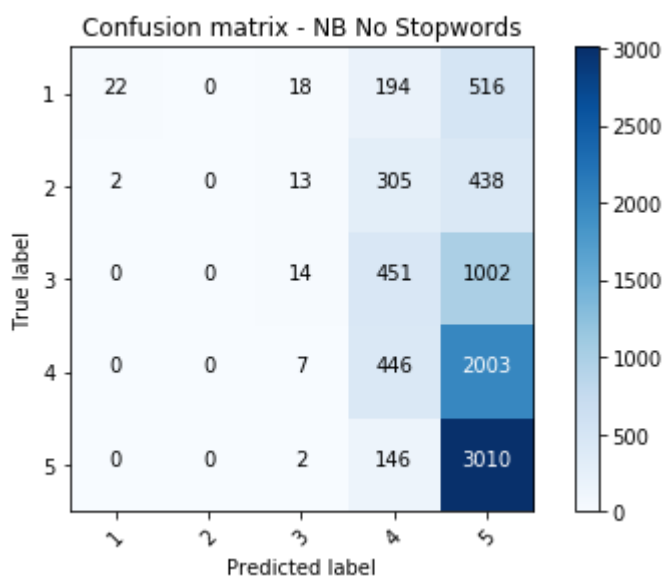

Confusion matrix, without normalization

```
[[ 428   84   86   61   91]
 [ 175  132  224  149   78]
 [   87  100  443  531  306]
 [   36   35  255 1091 1039]
 [   26   13   71  678 2370]]
```

Confusion matrix, without normalization

```
[[ 22   0   18  194  516]
 [   2   0   13  305  438]
 [   0   0   14  451 1002]
 [   0   0   7   446 2003]
 [   0   0   2   146 3010]]
```






```
In [40]: from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer

hotel_review_data['reviewsrating'] = hotel_review_data['reviewsrating']
hotel_review_data['reviewsrating'].value_counts()
```

```
Out[40]: 5    12671
         4     9880
         3     5747
         1     3070
         2     2988
         Name: reviewsrating, dtype: int64
```

```
In [41]: import matplotlib.pyplot as plt
fig = plt.figure(figsize=(8,6))
hotel_review_data['reviewsrating'].value_counts().plot.bar(ylim=0)
plt.show()
```



```
In [42]: reviewsrating_1 = hotel_review_data[hotel_review_data["reviewsrating"]==1]
reviewsrating_2 = hotel_review_data[hotel_review_data["reviewsrating"]==2]
reviewsrating_3 = hotel_review_data[hotel_review_data["reviewsrating"]==3]
reviewsrating_4 = hotel_review_data[hotel_review_data["reviewsrating"]==4]
reviewsrating_5 = hotel_review_data[hotel_review_data["reviewsrating"]==5]

subset_of_reviewsrating_1 = reviewsrating_1.sample(n=3000)
subset_of_reviewsrating_2 = reviewsrating_2.sample(n=2000)
subset_of_reviewsrating_3 = reviewsrating_3.sample(n=4000)
subset_of_reviewsrating_4 = reviewsrating_4.sample(n=5000)
subset_of_reviewsrating_5 = reviewsrating_5.sample(n=7000)

data_clean1 = pd.concat([subset_of_reviewsrating_1,subset_of_reviewsrating_2,subset_of_reviewsrating_3,subset_of_reviewsrating_4,subset_of_reviewsrating_5])
data_clean1 = data_clean1.sample(frac=1).reset_index(drop=True)
print("Current shape of dataset :",data_clean1.shape)
```

Current shape of dataset : (21000, 23)

```

In [43]: from sklearn.naive_bayes import MultinomialNB
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.linear_model import SGDClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

x = data_clean1.reviewstext
y = data_clean1.reviewsrating
count_vect = CountVectorizer()
x_train1, x_test1, y_train1, y_test1 = train_test_split(x, y, test_size=0.20, random_state=42)
X_train_counts = count_vect.fit_transform(x_train1)

tfidf_transformer = TfidfTransformer()
tf_transformer = TfidfTransformer(use_idf=False).fit(X_train_counts)
X_train_tf = tf_transformer.transform(X_train_counts)

X_train_tfidf = tfidf_transformer.fit_transform(X_train_counts)
clf = MultinomialNB().fit(X_train_tfidf, y_train1)

X_test_counts = count_vect.transform(x_test1)
tfidf_transformer = TfidfTransformer()
tf_transformer = TfidfTransformer(use_idf=False).fit(X_test_counts)
X_test_tf = tf_transformer.transform(X_test_counts)

pred = clf.predict(X_test_tf)
print("Multinomial Naive Bayes:", accuracy_score(y_test1, pred))

classif = SGDClassifier(loss='hinge', penalty='l2', alpha=1e-3, random_state=42, max_iter=1000)
clf1 = classif.fit(X_train_tfidf, y_train1)

pred = clf1.predict(X_test_tf)
print("SGD:", accuracy_score(y_test1, pred))

clf2 = DecisionTreeClassifier(random_state=0)
clf3 = clf2.fit(X_train_tfidf, y_train1)
pred = clf3.predict(X_test_tf)
print("Decision Tree:", accuracy_score(y_test1, pred))

clf4 = RandomForestClassifier(n_estimators=100, max_depth=5, random_state=0)
clf5 = clf4.fit(X_train_tfidf, y_train1)
pred = clf5.predict(X_test_tf)
print("Random Forest:", accuracy_score(y_test1, pred))

from sklearn.pipeline import Pipeline
text_clf = Pipeline([('vect', CountVectorizer()),
                     ('tfidf', TfidfTransformer()),
                     ('clf-svm', SGDClassifier(loss='hinge', penalty='l2', alpha=1e-3, max_iter=1000))])
text_clf = text_clf.fit(x_train1, y_train1)

predicted = text_clf.predict(x_test1)
np.mean(predicted == y_test1)

```

Multinomial Naive Bayes: 0.385476190476

SGD: 0.472142857143

Decision Tree: 0.360714285714

Random Forest: 0.340238095238

C:\Users\vijay\AppData\Local\conda\conda\envs\my_root\lib\site-packages\sklearn
\linear_model\stochastic_gradient.py:117: DeprecationWarning:

n_iter parameter is deprecated in 0.19 and will be removed in 0.21. Use max_ite
r and tol instead.

Out[43]: 0.49904761904761907


```

In [52]: data_clean_rating15 = pd.concat([subset_of_reviewsrating_1,subset_of_reviewsrating_2])
data_clean_rating15 = data_clean_rating15.sample(frac=1).reset_index(drop=True)
print("Current shape of dataset :",data_clean_rating15.shape)

y = data_clean_rating15.reviewsrating
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

#Setting up the X and Y data, where X is the review text and Y is the rating
#Three different inputs will be used: original review text, cleaned review text, and adjreview

x1 = data_clean_rating15.reviewstext
x2 = data_clean_rating15.cleanrev
x3 = data_clean_rating15.adjreview

x3.head()
#Creating a vectorizer to split the text into unigrams and bigrams
vect = TfidfVectorizer(ngram_range = (1,2))
x_vect1 = vect.fit_transform(x1)
x_vect2 = vect.fit_transform(x2)
x_vect3 = vect.fit_transform(x3)

#Making some simple functions for linear svc, knn, and naive bayes
def linsvc(x,y):
    x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.20, random_state=42)
    classf = LinearSVC()
    classf.fit(x_train, y_train)
    # pred1 = classf.predict(x_train)
    # print("Linear SVC: Training Set",accuracy_score(y_train, pred1))
    pred = classf.predict(x_test)
    print("Linear SVC:",accuracy_score(y_test, pred))
    return(y_test, pred)

def revknn(x,y):
    x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=42)
    classf = KNeighborsClassifier(n_neighbors=3)
    classf.fit(x_train, y_train)
    pred = classf.predict(x_test)
    print("kNN:",accuracy_score(y_test, pred))
    return(y_test, pred)

def revnb(x,y):
    x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=42)
    classf = MultinomialNB()
    classf.fit(x_train, y_train)
    pred = classf.predict(x_test)
    print("Naive Bayes:",accuracy_score(y_test, pred))
    return(y_test, pred)

```

Current shape of dataset : (10000, 23)

```
In [53]: from sklearn.svm import LinearSVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import MultinomialNB, GaussianNB

svmy1,svmp1 = linsvc(x_vect1,y)
svmy2,svmp2 = linsvc(x_vect2,y)
svmy3,svmp3 = linsvc(x_vect3,y)

knny1,knnp1 = revknn(x_vect1,y)
knny2,knnp2 = revknn(x_vect2,y)
knny3,knnp3 = revknn(x_vect3,y)

nby1,nbp1 = revnb(x_vect1,y)
nby2,nbp2 = revnb(x_vect2,y)
nby3,nbp3 = revnb(x_vect3,y)
```

```
Linear SVC: 0.93
Linear SVC: 0.929
Linear SVC: 0.878
kNN: 0.7176
kNN: 0.718
kNN: 0.7364
Naive Bayes: 0.8228
Naive Bayes: 0.8436
Naive Bayes: 0.8396
```

```
In [54]: print(clf.predict(count_vect.transform(["A great Hotel and perfect for a traveler
[5]
```

```
In [55]: print(clf.predict(count_vect.transform(["2 people checking in a long line Was bad
[1]
```

```
In [56]: print(clf.predict(count_vect.transform(["This hotel was nice and a great value fo
[5]
```

```
In [57]: print(clf.predict(count_vect.transform(["Bed bugs, will not be back"])))
[1]
```